

Reduce Performance Impact of Compaction Process by Designing a Global Format for LSM

Jinghuan YU
City University of Hong Kong

Chun Jason. Xue
City University of Hong Kong

Abstract

Your abstract text goes here. Just a few facts. Whet our appetites. Not more than 200 words, if possible, and preferably closer to 150.

1 Introduction

2 Background

Some Quick Conclusion HERE

2.1 Non-Volatile Memory

NVM, or PM (persistent memory), SCM (storage class memory), is actually the same meaning, referring to a series of memory materials with non-volatile and byte-wise access characteristics. NVM has become a hot topic in recent years, and related research is moving forward. There are many different types materials like PCM (phase change memory), MRAM (Magnetoresistive RAM) etc. Beside the most obvious feature, NVM has higher throughput and storage density than traditional NAND flash devices, with the shortages like asymmetrical read/write performance, and limited lifetime leads to the wear-out problems.

Although there is no final conclusion on how to use this material, research proposed several main solutions for using this material:

1. Use NVM as Persistent Transnational Memory, use methods like Redo Logging, Undo Logging, and Log-Structured to manage the transaction and data involved.
2. Use NVM as a disk to provide better random access performance on blocks and files. For example, introducing the Direct Access (DAX) in Linux. There are also file systems [11] similar to PMFS, NOVA [20], etc.
3. Combining with RDMA in a distributed scenario. Due to the high performance of NVM, the access characteristics of byte addressing, and the access mode of RDMA, distributed NVM + RDMA becomes a new architecture design.

2.2 Log-Structured-Merged Tree

LSM Tree is a high warm writing performance data structure proposed in 1996 [14], gets widely used in many products like Cassandra [1], Hbase [2], BigTable [9] and WiredTiger [6]. The most important characteristics LSM has are sequential write optimized and its periodic garbage collection.

2.2.1 Sequential Write Optimized

LSM's basic idea is converting random writes to sequential writes. Most of the storage devices has the characteristics that can perform much better in sequential operations than in random access operations, no matter read or write. LSM caches the newest changes inside the memory and use batch processing to write down those cached data. In addition to simply aggregate the operations, recent updated data will overwrite in memory and be flash to the file system only for the last version.

2.2.2 Periodic Garbage Collection

For using log-structure and incremental update to persistent operations. Data stored may be outdated periodically due to deletions and updates, so garbage collection process known as "Compaction" is needed. During the compaction process, LSM-based system will iterate through part of the persistent data, delete the out-dated data and reorganize the data in partial sorted manner to keep read optimal.

For this *Compaction* process will walk through all key-value pairs stored in the target files, apply merge sorting to these entries and write back to the files. Though most of implementations tried to optimize compaction process, there are still very severe performance impact and resource occupying problems during this process. Fig. 1 shows the disk usage over time. These problems have been studied from many years ago, there are three main directions to solve these problems.

- The first one is **scheduling**, bLSM [17] proposed a "Gear Scheduler" to dispersion pressure caused by compaction.




Figure 1: This chart shows the sudden increasing disk usage while system fall into compaction.

It inspired many following works and most of the work mentioned in this paper used bLSM’s results as base line while discussing scheduling problems.

- The second one is **algorithm optimization**, the most successful and representative one is PebblesDB in 2017 [15], use “Guard” to avoid repeated writing entries into files, easing the pressure by reducing write amplification. There are also works focus on special workloads like LSM-trie [19]. And this may be the most popular thoughts in practice, Facebook provides several different types of compaction in RocksDB [3], [10] while Cassandra also changed many times for its compaction strategy [4].
- The third one is about **takeing advantages of new storage materials**, for new materials has been developed lot recent years, how these new materials’, like Open-Channel SSD [8] and NVM, characteristics may benefit the LSM data structure becomes a attractive topic to develop. GearDB [21] considered reorganizing the entries choosing strategy to cover the garbage collection on HM-SMR disks. FlashKV [22] use Open-Channel SSDs to optimize the compaction process’s write amplification, achieved higher GC-efficiency. Novelsm [13] utilized the characteristics of NVM’s byte-addressing to achieve in-place update while SLM-DB [12] combined with persistent B-Tree to organize the files into one single level and use *select-merging* to get better performance in compaction.

2.3 Chances and Challenges

Programming for NVM is very different from traditional memory or disk programming, this section describes several main challenges while combing the NVM with the LSM structures.

2.3.1 Space Amplification

Fundamentally, this problem is unavoidable for any log-based system as the trade off for update costs and point looking




Figure 2: The introduce of space amplification, the different reasons’ ratio in entire result, point out most of the amplification problem is caused by the Compaction

up costs. Some hardware devices like HM-SMR and Open-channel may provide raw disk control to the applications, make it is possible for application to cover this problem with device’s own garbage collection process [22].

This problem can be the most typical shortage of LSM-based systems, PebblesDB [15] and RocksDB [10] tried to solve this problem by optimizing the write down process, while Novelsm [13] applied in-place updates on NVM to reduce repeated writing and SLM-DB use persistent B-Tree to solve this problem.

NVM’s byte addressing ability allows more flexible data structure and operations; Its large capacity also benefits the buffer to store much more data, caching more operations before writing down to the sequential-based devices. This can reduce space amplification from the very origin purpose.

2.3.2 Consistency

the most important is the consistency problem. For example, consider about the situation of double-linked table insert operation, while the following code executes to the second line and meets a sudden power failure.

```
void list_add_tail(struct cds_list_head *newp,
struct cds_list_head *head) {
head->prev->next = newp;
newp->next = head;
newp->prev = head->prev;
head->prev = newp;
}
```

For it is non-volatile in the NVM, when the system is restored, the linked list is in an abnormal state caused by the CPU Cache and the out-of-order execution of CPU. This means NVM requires a specified *transaction programming* model [11, 16] to ensure the semantics of atomic operations are achieved, which means there is no intermediate state generated during the system is restoring, this results in the extra memory fence and cache flush operations to keep consistency. Moreover, to provide consistent writing on data blocks larger




Figure 3: This chart shows the ratio of time used on reading file block from disk during compaction.

than limited size (8 bytes in typical situations), logging and C-o-W (Copy on Write) is needed.

2.3.3 Frequent serialization and de-serialization

LSM system achieved buffering data inside memory to provide append-only writing strategy, and many of implementations use *MemTable* as in-memory buffer and *SST* (Sorted String Table) files to store Key-Value pairs in the file system.

Use LevelDB [5] as an example, while applying a *Get* operation, LevelDB will first find in the memory buffer (*Memtable*), then access the version control to check which file contains the target key-value pair; After locating the file, it still needs use several iterator to transfer the file blocks to memory structure, and decode from pre-compression format to get the real value.

Fig. 3 shows a result that is not so common sense, the time spent on serialization and de-serialization take the most part of the system, much higher than the time spent on reading file blocks.

3 Proposed Solution

The challenge is to design a LSM-based system with only one single data structure which can

3.0.1

References

- [1] Apache cassandra. <http://cassandra.apache.org/>. (Accessed on 03/07/2019).
- [2] Apache hbase – apache hbase™ home. <https://hbase.apache.org/>. (Accessed on 03/07/2019).
- [3] Compaction · facebook/rocksdb wiki. <https://github.com/facebook/rocksdb/wiki/Compaction>. (Accessed on 03/10/2019).
- [4] Documentation. <http://cassandra.apache.org/doc/latest/operating/compaction.html?highlight=compaction>. (Accessed on 03/10/2019).
- [5] Leveldb.org. <http://leveldb.org/>. (Accessed on 03/07/2019).
- [6] Wiredtiger: making big data roar. <http://www.wiredtiger.com/>. (Accessed on 03/07/2019).
- [7] Remzi H. Arpaci-Dusseau and Arpaci-Dusseau Andrea C. *Operating Systems: Three Easy Pieces*. Arpaci-Dusseau Books, LLC, 1.00 edition, 2015. <http://pages.cs.wisc.edu/~remzi/OSTEP/>.
- [8] Matias Bjørling, Javier González, and Philippe Bonnet. Lightnvm: The linux open-channel {SSD} subsystem. In *15th {USENIX} Conference on File and Storage Technologies ({FAST} 17)*, pages 359–374, 2017.
- [9] Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C Hsieh, Deborah A Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E Gruber. Bigtable: A distributed storage system for structured data. *ACM Transactions on Computer Systems (TOCS)*, 26(2):4, 2008.
- [10] Siying Dong, Mark Callaghan, Leonidas Galanis, Dhruva Borthakur, Tony Savor, and Michael Strum. Optimizing space amplification in rocksdb. In *CIDR*, volume 3, page 3, 2017.
- [11] Subramanya R Dulloor, Sanjay Kumar, Anil Keshavamurthy, Philip Lantz, Dheeraj Reddy, Rajesh Sankaran, and Jeff Jackson. System software for persistent memory. In *Proceedings of the Ninth European Conference on Computer Systems*, page 15. ACM, 2014.
- [12] Olzhas Kaiyrakhmet, Songyi Lee, Beomseok Nam, Sam H Noh, and Young-ri Choi. Slm-db: Single-level key-value store with persistent memory. In *17th {USENIX} Conference on File and Storage Technologies ({FAST} 19)*, pages 191–205, 2019.
- [13] Sudarsun Kannan, Nitish Bhat, Ada Gavrilovska, Andrea Arpaci-Dusseau, and Remzi Arpaci-Dusseau. Redesigning lsms for nonvolatile memory with novelsm. In *2018 {USENIX} Annual Technical Conference ({USENIX}{ATC} 18)*, pages 993–1005, 2018.
- [14] Patrick O’Neil, Edward Cheng, Dieter Gawlick, and Elizabeth O’Neil. The log-structured merge-tree (lsm-tree). *Acta Informatica*, 33(4):351–385, 1996.
- [15] Pandian Raju, Rohan Kadekodi, Vijay Chidambaram, and Ittai Abraham. Pebblesdb: Building key-value stores using fragmented log-structured merge trees. In *Proceedings of the 26th Symposium on Operating Systems Principles*, pages 497–514. ACM, 2017.

- [16] Jinglei Ren, Jishen Zhao, Samira Khan, Jongmoo Choi, Yongwei Wu, and Onur Mutiu. Thynvm: Enabling software-transparent crash consistency in persistent memory systems. In *2015 48th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, pages 672–685. IEEE, 2015.
- [17] Russell Sears and Raghu Ramakrishnan. blsm: a general purpose log structured merge tree. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data*, pages 217–228. ACM, 2012.
- [18] Carl A. Waldspurger. Memory resource management in VMware ESX server. In *USENIX Symposium on Operating System Design and Implementation (OSDI)*, pages 181–194, 2002. <https://www.usenix.org/legacy/event/osdi02/tech/waldspurger/waldspurger.pdf>.
- [19] Xingbo Wu, Yuehai Xu, Zili Shao, and Song Jiang. Lsm-trie: An lsm-tree-based ultra-large key-value store for small data items. In *2015 {USENIX} Annual Technical Conference ({USENIX}{ATC} 15)*, pages 71–82, 2015.
- [20] Jian Xu and Steven Swanson. {NOVA}: A log-structured file system for hybrid volatile/non-volatile main memories. In *14th {USENIX} Conference on File and Storage Technologies ({FAST} 16)*, pages 323–338, 2016.
- [21] Ting Yao, Jiguang Wan, Ping Huang, Yiwen Zhang, Zhiwen Liu, Changsheng Xie, and Xubin He. Geardb: A gc-free key-value store on hm-smr drives with gear compaction. In *17th {USENIX} Conference on File and Storage Technologies ({FAST} 19)*, pages 159–171, 2019.
- [22] Jiacheng Zhang, Youyou Lu, Jiwu Shu, and Xiongjun Qin. Flashkv: Accelerating kv performance with open-channel ssds. *ACM Transactions on Embedded Computing Systems (TECS)*, 16(5s):139, 2017.